

Introduction to XGboost

FU Hanlin

`fixed-term.hanlin.fu@bosch.com`

AE/MFP2.1-CN

April 10, 2019

Outline

- 1. Unsupervised Learning**
- 2. Supervised Learning**
- 3. Ensemble Learning**
- 4. XGboost**

1. Unsupervised Learning

Unsupervised Learning





Figure: Unsupervised Learning(e.g., Mickey Mouse and Donald Duck)

Clustering




→ Market segmentation



→ Social network analysis



→ Organize computing clusters



→ Astronomical data analysis

Figure: Clustering examples

K-means

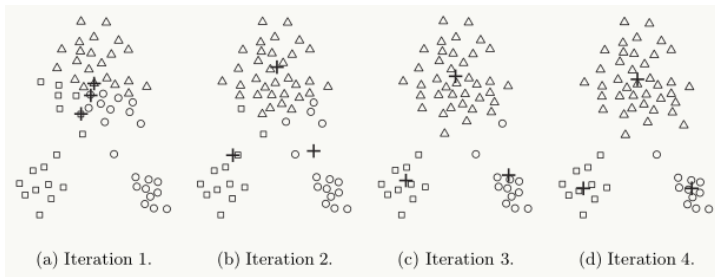


Figure: K-means Iteration(3 clusters)

K-means algorithm

Input:

- K (number of clusters)
- Training set $x_{(1)}, x_{(2)}, \odot, x_{(m)}$

2. Supervised Learning

Classification




Figure: A labeled training set for supervised learning (e.g., spam classification)

Regression




Figure: Regression

Classification Tree

Classification tree analysis is when the predicted outcome is the class (discrete) to which the data belongs.

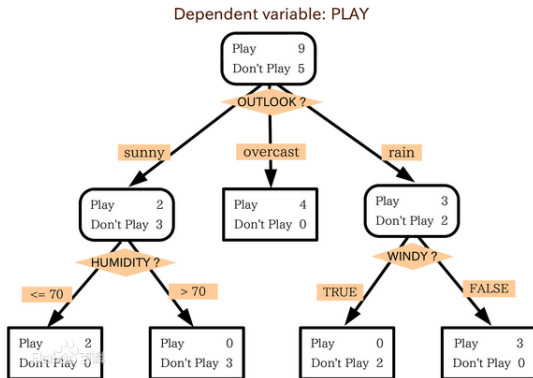


Figure: A decision tree with binary splits for classification

How to split

generative algorithm	categorization
<i>ID3</i>	information entropy
<i>C4.5</i>	gain ratio
<i>CART</i>	Gini index


Regression Tree(CART)

Regression tree analysis is when the predicted outcome can be considered a real number.

- regression tree (also known as classification and regression tree)
- contains one score in each leaf value
- base learner

Input: age, gender, occupation, ...

Does the person like computer games



prediction score in each leaf

+2

+0.1

-1

How to predict

Step 1: Choose split point $\sum (y_i - f(x_i))^2$

- Area1: $R_1(j, s) = \{x | x^{(j)} \leq s\}$
- Area2: $R_2(j, s) = \{x | x^{(j)} > s\}$

Step 2: Minimize loss function

- $\min_{c_1} \sum_{x_i \in R_1(j, s)} (y_i - c_1)^2 + \min_{c_2} \sum_{x_i \in R_2(j, s)} (y_i - c_2)^2$

Step 3: Recursively call Step 1 and Step 2 in subareas

Step 4: Output predicted value

- $c_1 = \text{ave}(y_i | x_i \in R_1(j, s))$
- $c_2 = \text{ave}(y_i | x_i \in R_2(j, s))$

3. Ensemble Learning

Ensemble Theory

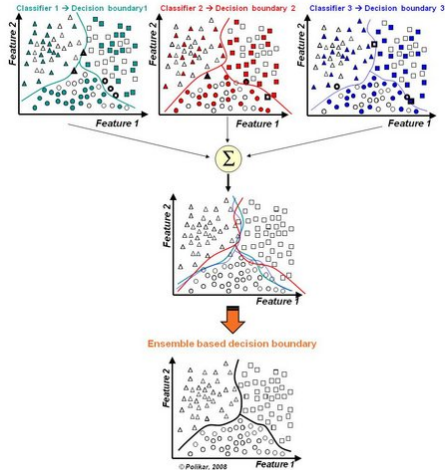


Figure: Model stacking

Model Combination

\sum in previous slide

- averaging

1. simple averaging: $H(x) = \frac{1}{T} \sum_{i=1}^T h_i(x)$

2. weighted averaging: $H(x) = \sum_{i=1}^T \omega_i h_i(x)$

- voting

1. majority voting

2. plurality voting

3. weighted voting

- stacking

Bagging and Boosting




Figure: Bagging




Figure: Bootstrap Aggregating

4. XGboost

eXtreme Gradient Boosting Objective

$$Obj(\Theta) = L(\Theta) + \Omega(\Theta)$$

$L(\Theta)$ Training Loss, measures how well model fit on training data

$\Omega(\Theta)$ Regularization, measures complexity of model

- Optimizing **training loss** encourages **predictive** models
Fitting well in training data at least get you close to training data which is hopefully close to the underlying distribution.
- Optimizing **regularization** encourages **simple** models
Simpler models tends to have smaller variance in future predictions, making prediction **stable**

Learning Step Function(visually)




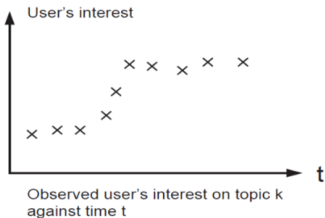


Figure: $L(f)$ VS $\Omega(f)$

How do we learn?

- Objective: $\sum_{i=1}^n l(\hat{y}_i, y) + \Omega(f_k), \Omega(f_k) = \gamma T + \frac{1}{2} \lambda ||w||^2, f_k \in F$

- Solution: Boosting

Start from constant prediction, add a new function each time

$$\hat{y}_i^{(0)} = 0$$

$$\hat{y}_i^{(1)} = f_1(x_i) = \hat{y}_i^{(0)} + f_1(x_i)$$

$$\hat{y}_i^{(2)} = f_1(x_i) + f_2(x_i) = \hat{y}_i^{(1)} + f_2(x_i)$$

...

$$\hat{y}_i^{(t)} = \sum_{k=1}^t f_k(x_i) = \hat{y}_i^{(t-1)} + f_t(x_i)$$

- Additive training

- Structure score: $Obj = -\frac{1}{2} \sum_{j=1}^T \frac{G_j^2}{H_j + \lambda} + \gamma T$

$G_j = \sum_{i \in I_j} \partial_{\hat{y}(t-1)} l(y_i, \hat{y}(t-1)), H_j = \sum_{i \in I_j} \partial_{\hat{y}(t-1)}^2 l(y_i, \hat{y}(t-1)), T$ is number of leaves, w are leaf scores

The Stucture Score Calculation

Instance index gradient statistics

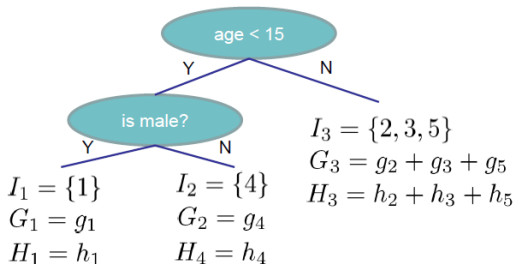
1  g1, h1

2  g2, h2

3  g3, h3

4  g4, h4


5  g5, h5



$$Obj = -\sum_j \frac{G_j^2}{H_j + \lambda} + 3\gamma$$

The smaller the score is, the better the structure is

How to split?



In practice, we grow the tree greedily

- Start from tree with depth 0
- For each leaf node of the tree, try to add a split. The change of objective after adding the split is

$$Gain = \frac{1}{2} \left[\frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda} \right] - \gamma$$

γ means the complexity cost by introducing additional leaf

References I

- [1] Tianqi Chen's Homepage
<https://homes.cs.washington.edu/~tqchen/>
- [2] Chen and C Gusterin. XGBoost: A Scalable Tree Boosting System. In KDD 16—Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, page 785-794.
- [3] Apache
<http://xgboost.apachecn.org/cn/latest/model.html>